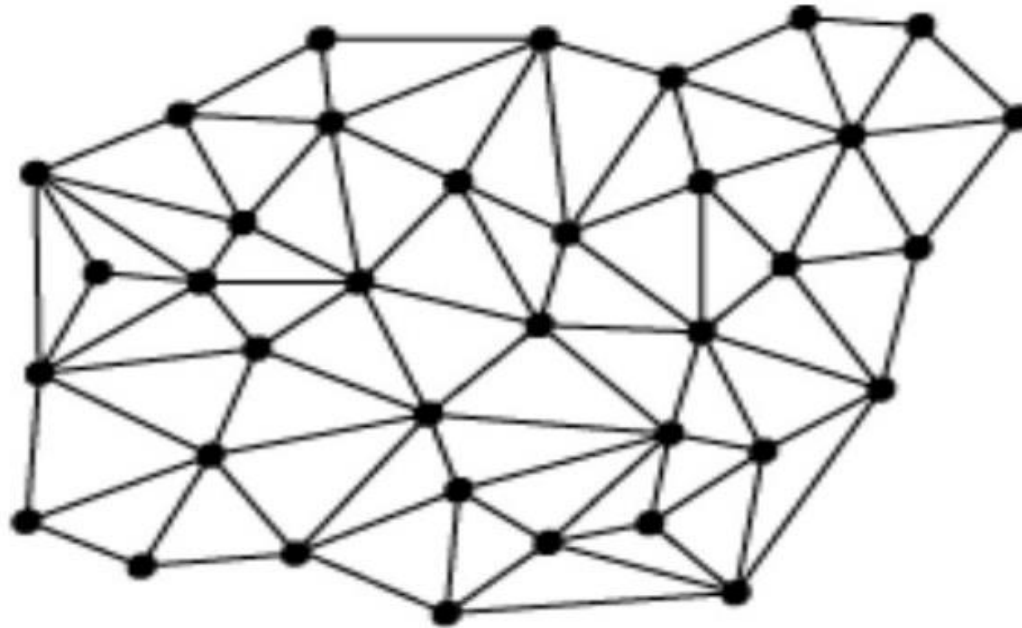


# Distributed System Course

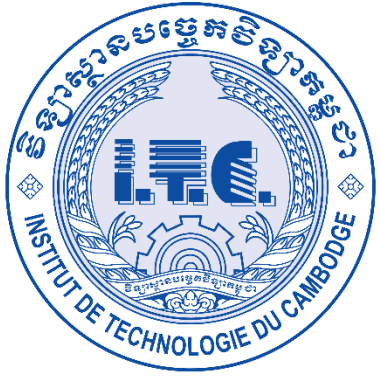
2021-22-GICI41SSD-Distributed System



Academic Year: 2021-2022    Lecturer: SOK Kimheng

# Information

Course	Distributed System	48h, 12 Weeks, 4h/week (3 Groups = 96h)
General Distributed System	Week 1	Information, Self-Study Skill, Introduction
	Week 2	Distributed Communication (TCP/IP, Socket, RPC, REST, gRPC, OMQ)
	Week 3	Clock, Timestamp
	Week 4	Fault Tolerance (Two general problem, Byzantine General Problem)
	Week 5	Consensus Algorithm (Paxos, ZooKeeper, Raft)
	Week 6	Quiz
Blockchain	Week 7	Basic Cryptography
	Week 8	Blockchain and Bitcoin (Proof of Work)
	Week 9	Ethereum and Smart Contract (Proof of Stake)
	Week 10	Hyperledger and Self-Sovereign Identity
	Week 11	Security
	Week 12	Final Exam



# Distributed System Course

2021-22-GICI41SSD-Distributed System

Week5:

## Consensus Algorithm

Paxos, Zookeepr, Raft

Academic Year: 2021-2022    Lecturer: SOK Kimheng

# Agenda

- 1 Definition
- 2 Case study
- 3 Paxos
- 4 ZooKeeper
- 5 Raft

# Consensus Algorithm

## Definition

- A general agreement
- A collectively agreement on the same output by all nodes | peers within the distributed system

# Consensus Algorithm

## Case study

1. A group of students decide to make a group T-shirt
2. Choosing T-shirt color.
  - ❖ Let's see who is initiator, proposer, leader
3. Finding members, Creating rule
  - ❖ Split the group (divide, fork) , or unite by majority
4. Asking confirmation
  - ❖ Number of confirmation received
5. Commit the decision
  - ❖ What is the final color?
6. Leader term (Communism or Democracy?)

# Consensus Algorithm


## Case study

Why do systems need to reach consensus?




So, either

Leader-replicas schema



or

Peer-to-peer schema



# Consensus Algorithm

## PAXOS (1989, 2001)

- Consensus is agreement on one result
- Once a majority agrees on a proposal, that is the consensus
- The reached consensus can be eventually known by everyone
- Paxos defines 3 roles: Proposers, Acceptors, and Learners
- Paxos nodes can take multiple roles, even all of them
- Paxos nodes must know how many acceptors a majority is
- Paxos nodes must be persistent: they can't forget what they accepted



# Consensus Algorithm

## PAXOS (1989, 2001)

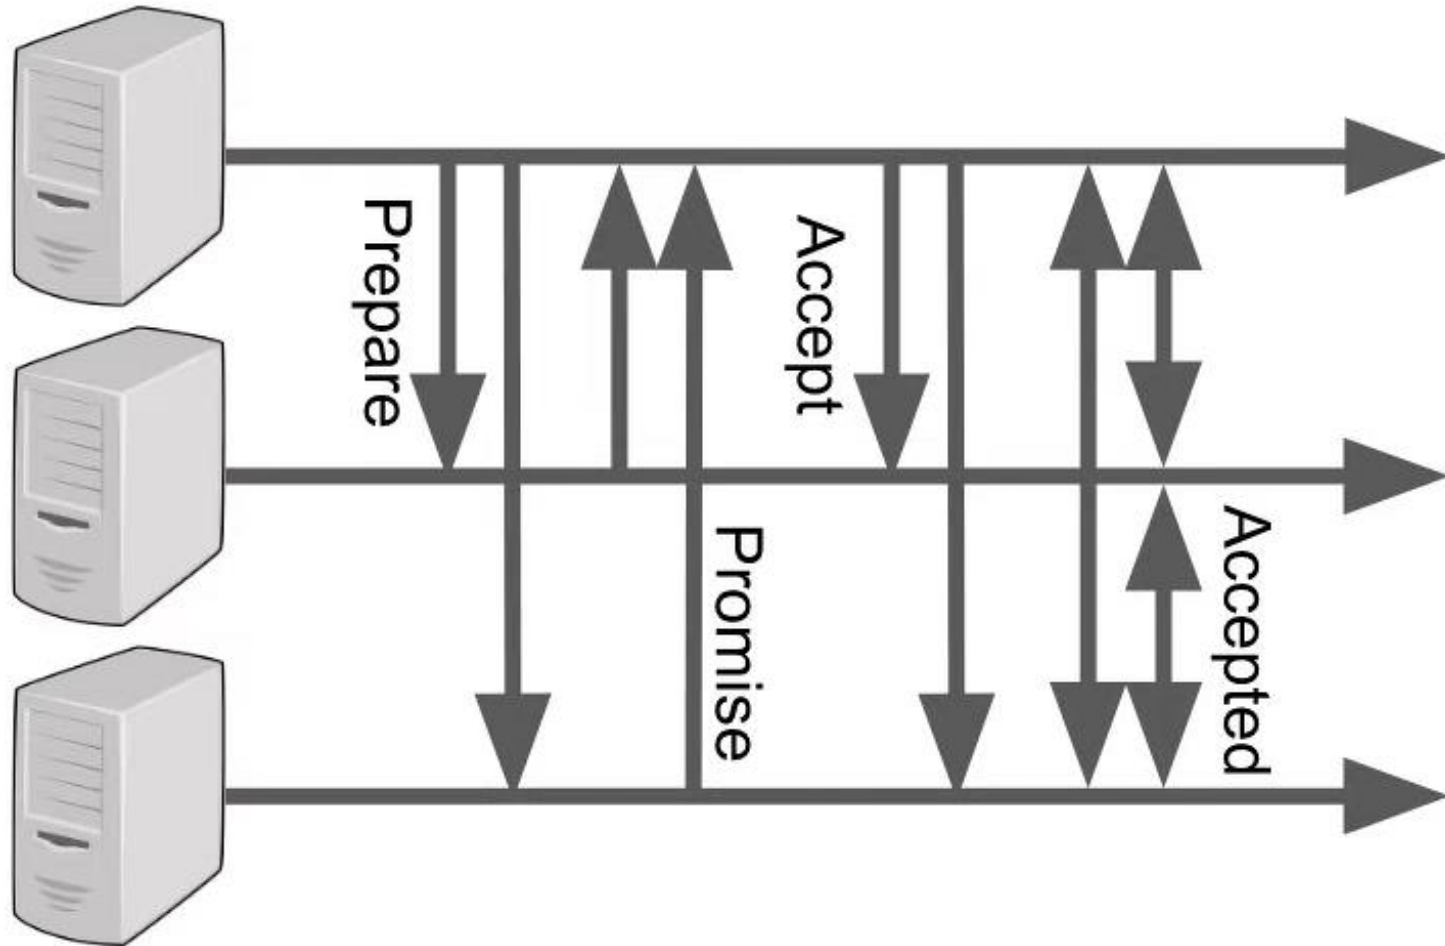
### Alternatives to Paxos



<b>Fault Types:</b>	Byzantine	Fail Stop	
<b>Failover:</b>	Instant		Takes time
<b>Servers:</b>	$3m+1$	$2m+1$	$m+1$
<b>Messages:</b>	Exponential	Linear	


# Consensus Algorithm

PAXOS (1989, 2001)



# Consensus Algorithm

## PAXOS (1989, 2001)




**fault**

- One node fail to send back promise
- Count the majority

# Consensus Algorithm

## PAXOS (1989, 2001)




**fault**

- One node fail to send back Accepted
- Count the majority

# Consensus Algorithm

## PAXOS (1989, 2001)

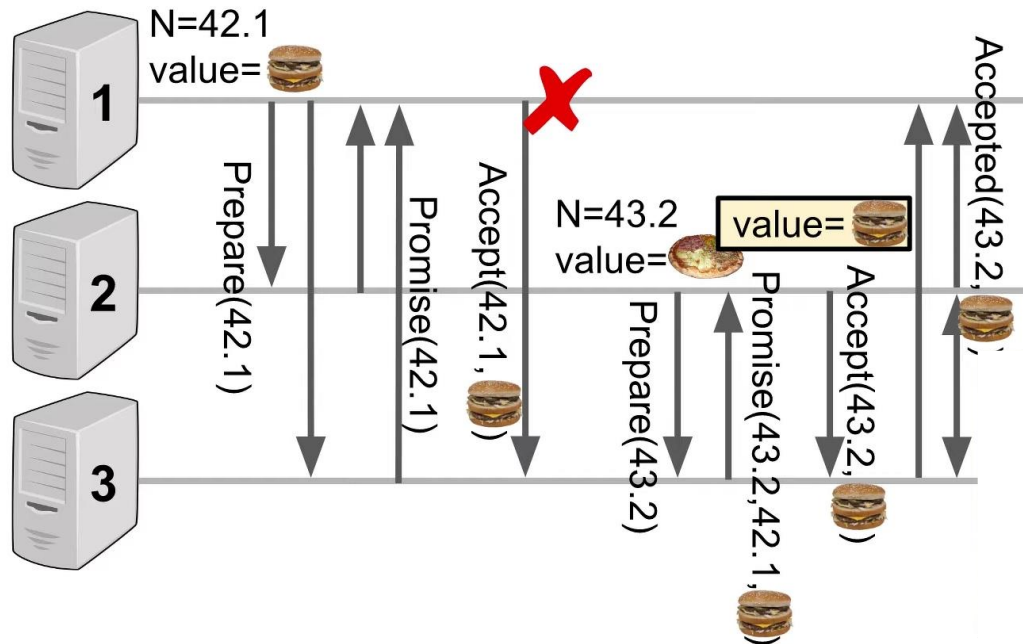


### fault

- First node fail after Prepare
- Second node make new prepare
- New decision is accepted

# Consensus Algorithm

## PAXOS (1989, 2001)



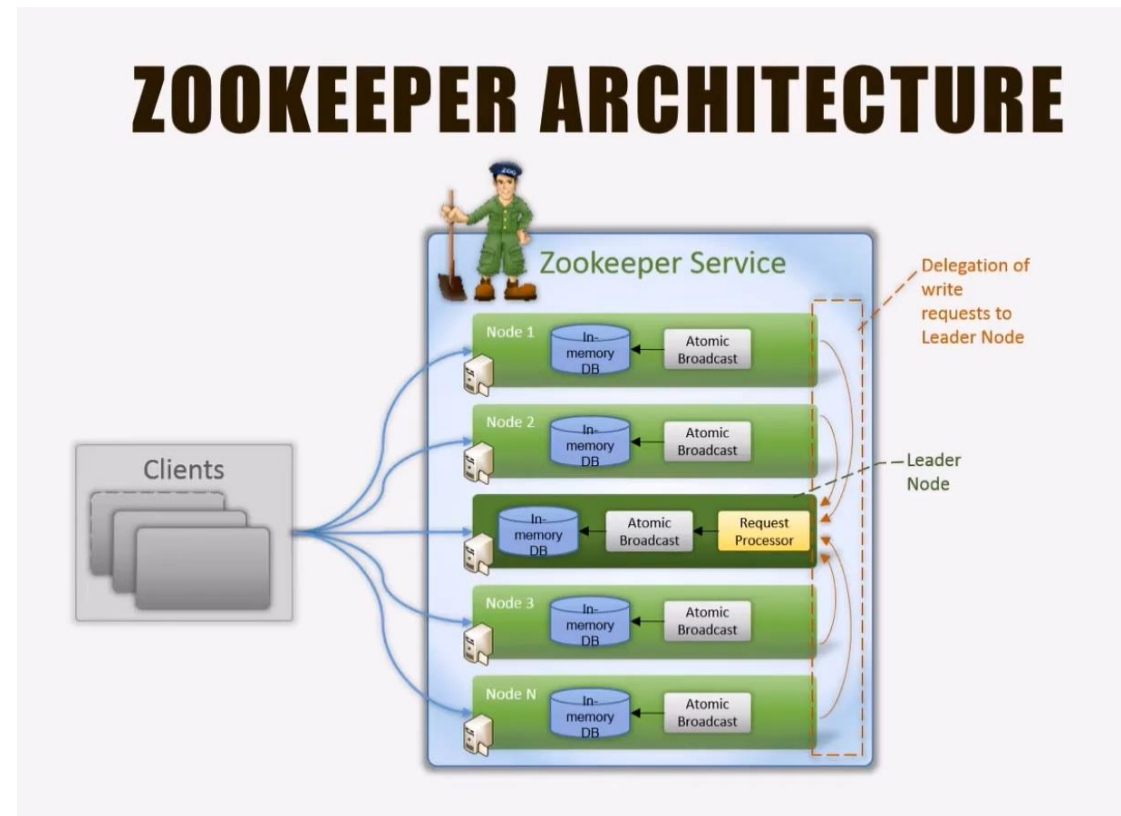
### fault

- First node fail during Accept
- Second node make new prepare
- Third node inform second node about previous value
- Second and Third nodes agree for the previous value

# Consensus Algorithm

## ZooKeeper

- Distributed and Open-source coordination service for decentralize applications

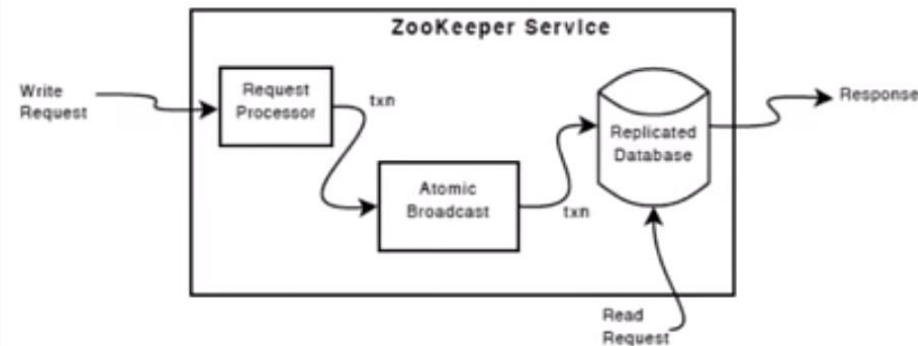


# Consensus Algorithm

## ZooKeeper

### IMPORTANT COMPONENTS

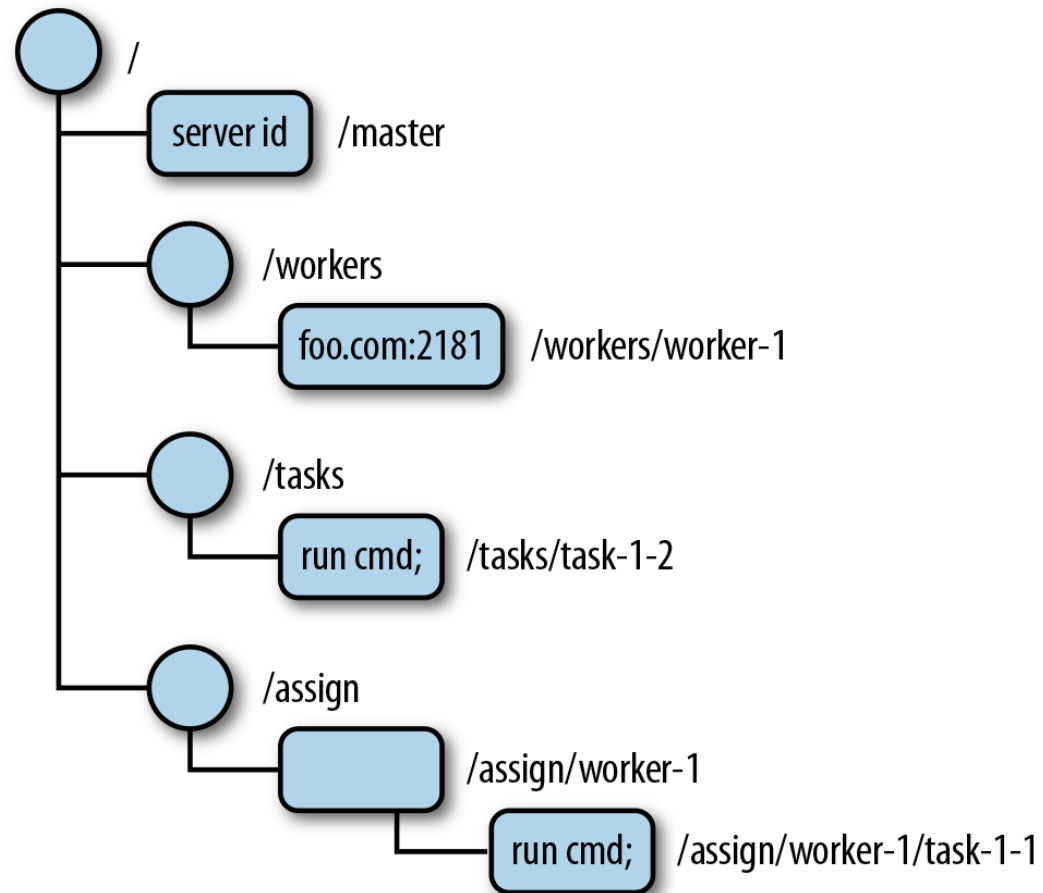
- **Leader & Follower**
- **Request Processor**
  - Active in Leader Node and is responsible for processing write requests.
  - After processing, it send changes to follower nodes
- **Atomic Broadcast**
  - Present in both Leader Node and Follower Nodes.
  - It is responsible for sending the changes to other nodes
- **In-memory Database (Replicated Database)**
  - It is responsible for storing the data in ZooKeeper.
  - Every node contains its own database
  - Data is also written to file system providing recoverability in case of any problems with cluster





# Consensus Algorithm

## ZooKeeper file structure

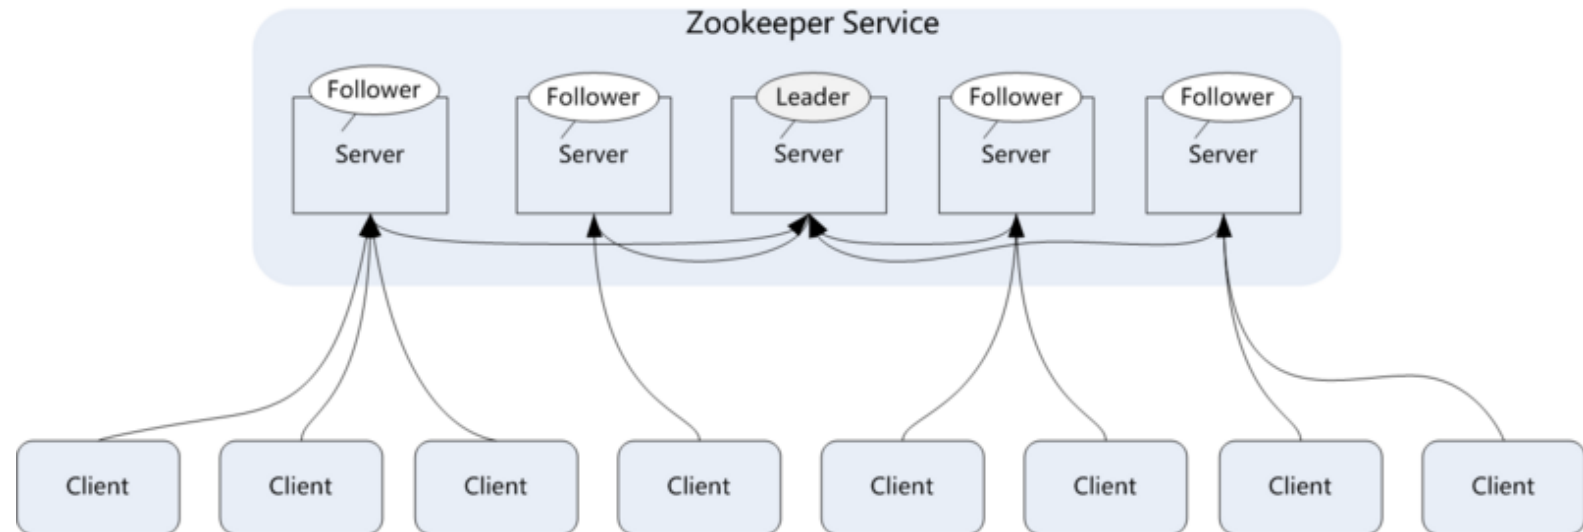


# Consensus Algorithm

## ZooKeeper

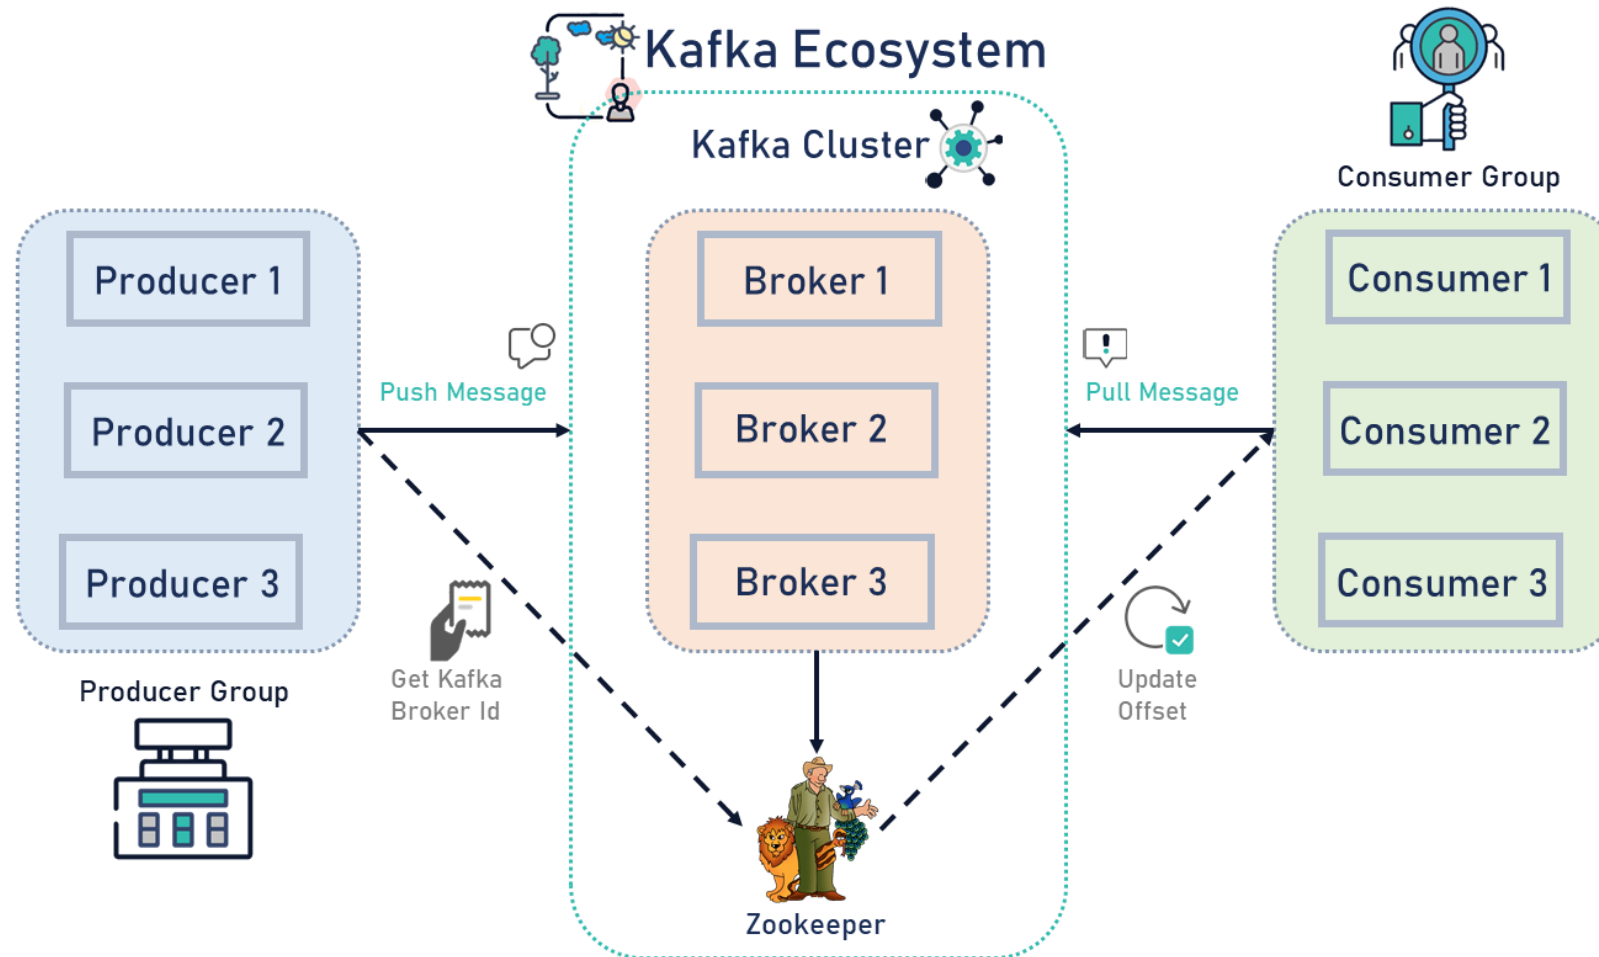
### ➤ ZooKeeper Maintains

- Version number
- ACL
- Time Stamp
- Data Length



# Consensus Algorithm

## ZooKeeper Problem



# Consensus Algorithm

## **Raft (Reliable | Replicate | Redundant and Fault Tolerant)**

### ➤ Leader Election

- Select one server to act as leader
- Detect crashes, choose new leader

### ➤ Log Replication

- Leader accepts commands from clients, appends to its log
- Leader replicates its log to other servers (overwrite if inconsistency)

### ➤ Safety

- Keep logs consistent
- Only servers with up-to-date logs can become leader

# Consensus Algorithm

## **Raft (Reliable | Replicate | Redundant and Fault Tolerant)**

### ➤ Roles / States

- Follower, Candidate, Leader

### ➤ Election Term

- At most 1 leader per term
- Some term has no leader (Failed election)
- Each server maintains current term value (No global view)
- Exchange communication through RPC
- If Leader encounters latest term, it updates its term and revert to follower
- If incoming RPC has obsolete term, reply with error

# Consensus Algorithm

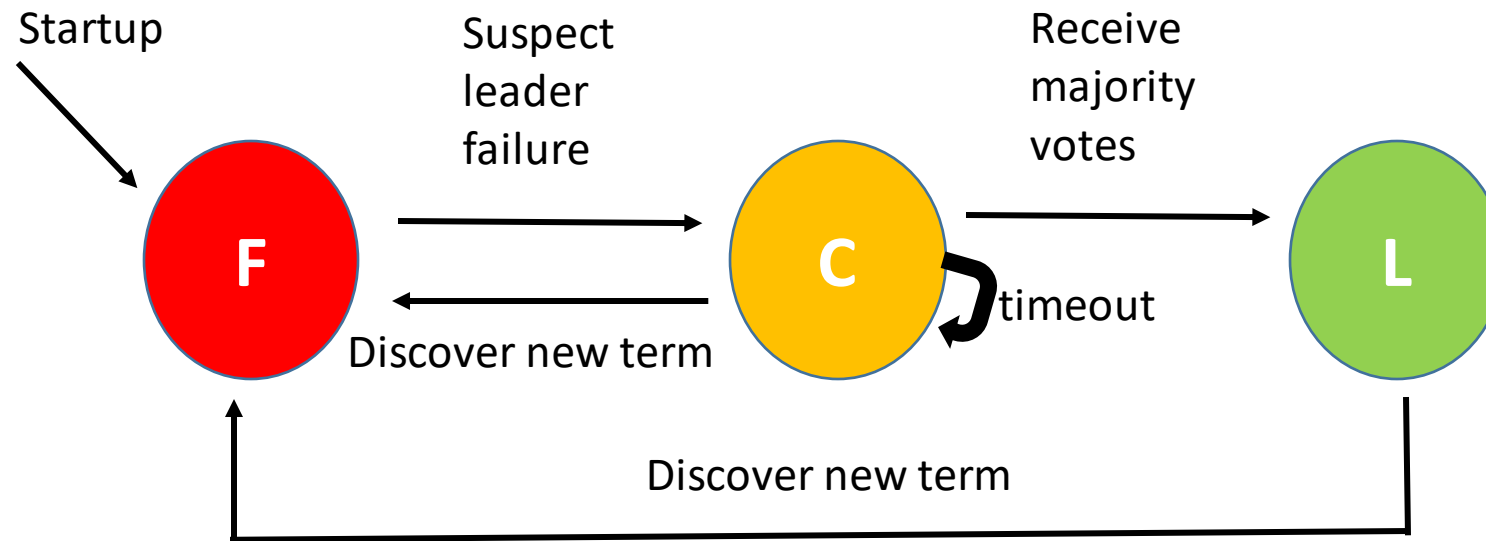
## **Raft (Reliable | Replicate | Redundant and Fault Tolerant)**

### ➤ Election Process

- Every node start with follower state
- Each nodes initiate a random election timeout
- If timeout, node changes from follower to candidate state, and start asking for vote
- If candidate received majority votes, it changes from candidate to leader state and start sending heartbeat to other nodes
- Leader will update the latest log to all the followers


# Consensus Algorithm

**Raft (Reliable | Replicate | Redundant and Fault Tolerant)**



# Consensus Algorithm

## Raft (Reliable | Replicate | Redundant and Fault Tolerant)




Suppose there are 5 nodes S1 to S5 in term 1 with random election time out.



# Consensus Algorithm


## Raft (Reliable | Replicate | Redundant and Fault Tolerant)



S2 timeout and start sending vote request.

# Consensus Algorithm


## Raft (Reliable | Replicate | Redundant and Fault Tolerant)



S2 received majority vote and becomes leader, then start sending heartbeat message to all nodes.

# Consensus Algorithm


**Raft (Reliable | Replicate | Redundant and Fault Tolerant)**



In case S3 and S4 are unresponsive, what will happen?

# Consensus Algorithm


## Raft (Reliable | Replicate | Redundant and Fault Tolerant)



S1 and S2 got only one vote from S5 and itself, so both of them can't become leader. So term 2 election is failed.

# Consensus Algorithm


## Raft (Reliable | Replicate | Redundant and Fault Tolerant)



When term 2 election timeout, S2 start term 3 election.

# Consensus Algorithm

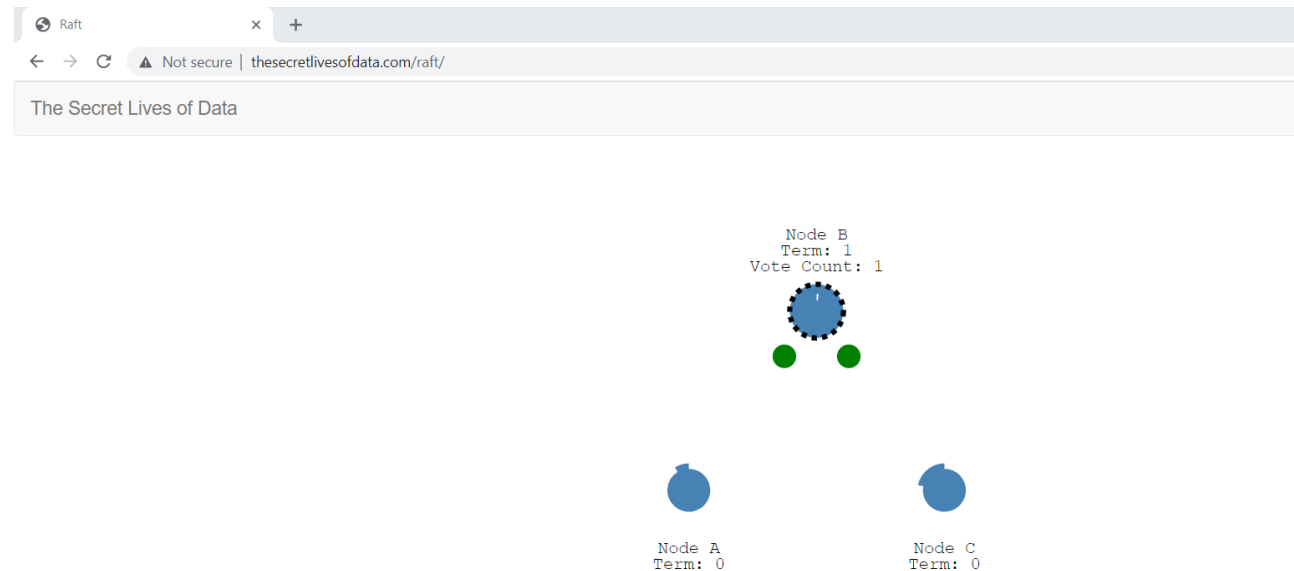
## Raft (Reliable | Replicate | Redundant and Fault Tolerant)



When S2 received majority vote, it becomes leader of term 3 and start sending heartbeat message to all nodes.

# Consensus Algorithm

## Raft with animation



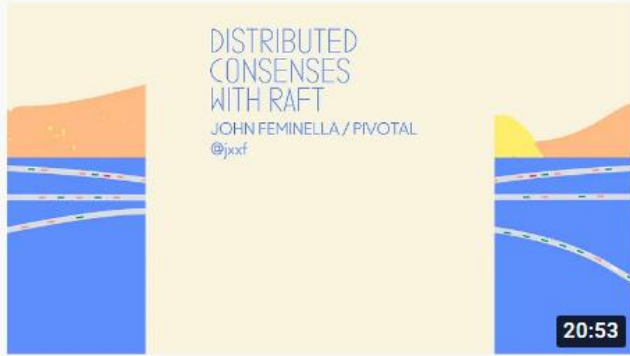
...and sends out *Request Vote* messages to other nodes.

Continue ➡

<http://thesecretlivesofdata.com/raft/>


# Consensus Algorithm

## Reference Videos

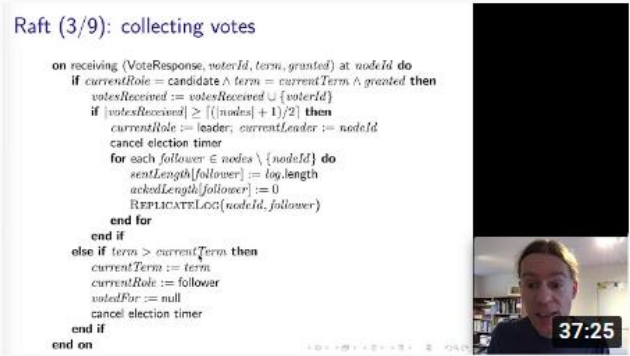


### Distributed Consensus with Raft - CodeConf 2016

16K views • 4 years ago

 GitHub

Presented by John Feminella Getting people to agree to things is sometimes hard. But implementing **consensus** w



### Raft (3/9): collecting votes

```
on receiving (VoteResponse, voterId, term, granted) at nodeId do
  if currentRole = candidate ^ term = currentTerm ^ granted then
    votesReceived := votesReceived ∪ {voterId}
    if |votesReceived| ≥ ((|nodes| + 1)/2) then
      currentRole := leader; currentLeader := nodeId
      cancel election timer
      for each follower ∈ nodes \ {nodeId} do
        sentLength[follower] := log.length
        ackedLength[follower] := 0
        REPLICATELOG(nodeId, follower)
      end for
    end if
  else if term > currentTerm then
    currentTerm := term
    currentRole := follower
    votedFor := null
    cancel election timer
  end if
end on
```

NOTE: There are some mistakes in this video. Please watch this one instead, in which the bugs are fixed: ...